

Cross Validation & Ensembling

Shan-Hung Wu
shwu@cs.nthu.edu.tw

Department of Computer Science,
National Tsing Hua University, Taiwan

Machine Learning

Outline

- 1 **Cross Validation**
 - How Many Folds?
- 2 **Ensemble Methods**
 - Voting
 - Bagging
 - Boosting
 - Why AdaBoost Works?

Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Cross Validation

- So far, we use the *hold out* method for:
 - Hyperparameter tuning: validation set
 - Performance reporting: testing set
- What if we get an “unfortunate” split?

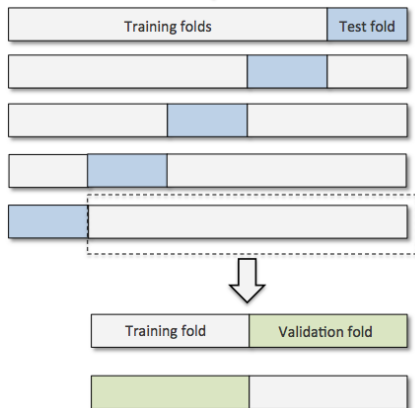
Cross Validation

- So far, we use the **hold out** method for:
 - Hyperparameter tuning: validation set
 - Performance reporting: testing set
- What if we get an “unfortunate” split?
- ***K-fold cross validation***:
 - ① Split the data set \mathbb{X} evenly into K subsets $\mathbb{X}^{(i)}$ (called ***folds***)
 - ② For $i = 1, \dots, K$, train $f_{-N^{(i)}}$ using all data but the i -th fold ($\mathbb{X} \setminus \mathbb{X}^{(i)}$)
 - ③ Report the ***cross-validation error*** C_{CV} by averaging all testing errors $C[f_{-N^{(i)}}]$'s on $\mathbb{X}^{(i)}$



Nested Cross Validation


- Cross validation (CV) can be applied to *both* hyperparameter tuning and performance reporting



- E.g, 5×2 *nested CV*

Nested Cross Validation


- Cross validation (CV) can be applied to *both* hyperparameter tuning and performance reporting



- E.g, 5×2 *nested CV*
- ① Inner (2 folds): select hyperparameters giving lowest C_{CV}
 - Can be wrapped by grid search

Nested Cross Validation


- Cross validation (CV) can be applied to **both** hyperparameter tuning and performance reporting



- E.g, 5×2 **nested CV**
- ① Inner (2 folds): select hyperparameters giving lowest C_{CV}
 - Can be wrapped by grid search
- ② Train final model using **both** training and validation sets with the selected hyperparameters

Nested Cross Validation

- Cross validation (CV) can be applied to **both** hyperparameter tuning and performance reporting



- E.g, 5×2 **nested CV**
- ① Inner (2 folds): select hyperparameters giving lowest C_{CV}
 - Can be wrapped by grid search
- ② Train final model using **both** training and validation sets with the selected hyperparameters
- ③ Outer (5 folds): report C_{CV} as test error

Outline

- 1 **Cross Validation**
 - How Many Folds?
- 2 **Ensemble Methods**
 - Voting
 - Bagging
 - Boosting
 - Why AdaBoost Works?

How Many Folds K ? I

How Many Folds K ? I

- The cross-validation error C_{CV} is an average of $C[f_{-N^{(i)}}]$'s

How Many Folds K ? I

- The cross-validation error C_{CV} is an average of $C[f_{-N^{(i)}}]$'s
- Regard each $C[f_{-N^{(i)}}]$ as an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$

How Many Folds K ? I

- The cross-validation error C_{CV} is an average of $C[f_{-N^{(i)}}]$'s
- Regard each $C[f_{-N^{(i)}}]$ as an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$
- C_{CV} is an estimator too, and we have

$$\text{MSE}(C_{CV}) = E_{\mathbb{X}}[(C_{CV} - E_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- *Mean square error* of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbf{x}} [(\hat{\theta}_n - \theta)^2]$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- Mean square error* of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2]$$

- Can be decomposed into the bias and variance:

$$\mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2] = \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2]$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- Mean square error** of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2]$$

- Can be decomposed into the bias and variance:

$$\begin{aligned}\mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2] &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2] \\ &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2 + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta)]\end{aligned}$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- Mean square error** of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2]$$

- Can be decomposed into the bias and variance:

$$\begin{aligned}\mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2] &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2] \\ &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2 + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta)] \\ &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + \mathbb{E} [(\mathbb{E}[\hat{\theta}_n] - \theta)^2] + 2\mathbb{E} (\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]) (\mathbb{E}[\hat{\theta}_n] - \theta)\end{aligned}$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- Mean square error** of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2]$$

- Can be decomposed into the bias and variance:

$$\begin{aligned}\mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2] &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2] \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2 + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta)] \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + \mathbb{E} [(\mathbb{E}[\hat{\theta}_n] - \theta)^2] + 2\mathbb{E} (\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]) (\mathbb{E}[\hat{\theta}_n] - \theta) \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2 \cdot 0 \cdot (\mathbb{E}[\hat{\theta}_n] - \theta)\end{aligned}$$

Point Estimation Revisited: Mean Square Error

- Let $\hat{\theta}_n$ be an estimator of quantity θ related to random variable \mathbf{x} mapped from n i.i.d samples of \mathbf{x}
- Mean square error** of $\hat{\theta}_n$:

$$\text{MSE}(\hat{\theta}_n) = \mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2]$$

- Can be decomposed into the bias and variance:


$$\begin{aligned}\mathbb{E}_{\mathbb{X}} [(\hat{\theta}_n - \theta)^2] &= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2] \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2 + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta)] \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + \mathbb{E} [(\mathbb{E}[\hat{\theta}_n] - \theta)^2] + 2\mathbb{E} (\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]) (\mathbb{E}[\hat{\theta}_n] - \theta) \\&= \mathbb{E} [(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2] + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2 \cdot 0 \cdot (\mathbb{E}[\hat{\theta}_n] - \theta) \\&= \text{Var}_{\mathbb{X}}(\hat{\theta}_n) + \text{bias}(\hat{\theta}_n)^2\end{aligned}$$

- MSE of an unbiased estimator is its variance

我們喜歡unbiased的 θ
因為它的MSE就是variance

Example: 5-Fold vs. 10-Fold CV


$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$



Example: 5-Fold vs. 10-Fold CV

$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$


- Consider polynomial regression where $P(y|x) = \sin(x) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma^2)$



Example: 5-Fold vs. 10-Fold CV

$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$


- Consider polynomial regression where $P(y|x) = \sin(x) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma^2)$
- Let $C[\cdot]$ be the MSE of predictions (made by a function) to true labels



Example: 5-Fold vs. 10-Fold CV

$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$


- Consider polynomial regression where $P(y|x) = \sin(x) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma^2)$
- Let $C[\cdot]$ be the MSE of predictions (made by a function) to true labels
- $\mathbb{E}_{\mathbb{X}}(C[f_N])$: read line



Example: 5-Fold vs. 10-Fold CV

$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$

- Consider polynomial regression where $P(y|x) = \sin(x) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma^2)$
- Let $C[\cdot]$ be the MSE of predictions (made by a function) to true labels
- $\mathbb{E}_{\mathbb{X}}(C[f_N])$: read line
- $\text{bias}(C_{\text{CV}})$: gaps between the red and other solid lines ($\mathbb{E}_{\mathbb{X}}[C_{\text{CV}}]$)




Example: 5-Fold vs. 10-Fold CV

$$\text{MSE}(C_{\text{CV}}) = \mathbb{E}_{\mathbb{X}}[(C_{\text{CV}} - \mathbb{E}_{\mathbb{X}}(C[f_N]))^2] = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2$$

- Consider polynomial regression where $P(y|x) = \sin(x) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma^2)$
- Let $C[\cdot]$ be the MSE of predictions (made by a function) to true labels
- $\mathbb{E}_{\mathbb{X}}(C[f_N])$: read line
- $\text{bias}(C_{\text{CV}})$: gaps between the red and other solid lines ($\mathbb{E}_{\mathbb{X}}[C_{\text{CV}}]$)
- $\text{Var}_{\mathbb{X}}(C_{\text{CV}})$: shaded areas

fold 較少 bias 較大 但 var 會較小



Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\text{bias}(C_{CV}) = E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N])$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N])\end{aligned}$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N]) \\ &= E(C[f_{-N^{(s)}}]) - E(C[f_N]), \forall s\end{aligned}$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N]) \\ &= E(C[f_{-N^{(s)}}]) - E(C[f_N]), \forall s \\ &= \text{bias}(C[f_{-N^{(s)}}]), \forall s\end{aligned}$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N]) \\ &= E(C[f_{-N^{(s)}}]) - E(C[f_N]), \forall s \\ &= \text{bias}(C[f_{-N^{(s)}}]), \forall s\end{aligned}$$

$$\text{Var}_{\mathbb{X}}(C_{CV}) = \text{Var}\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) = \frac{1}{K^2} \text{Var}\left(\sum_i C[f_{-N^{(i)}}]\right)$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$MSE(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N]) \\ &= E(C[f_{-N^{(s)}}]) - E(C[f_N]), \forall s \\ &= \text{bias}(C[f_{-N^{(s)}}]), \forall s\end{aligned}$$

$$\begin{aligned}\text{Var}_{\mathbb{X}}(C_{CV}) &= \text{Var}\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) = \frac{1}{K^2} \text{Var}\left(\sum_i C[f_{-N^{(i)}}]\right) \\ &= \frac{1}{K^2} \left(\sum_i \text{Var}(C[f_{-N^{(i)}}]) + 2 \sum_{i,j>i} \text{Cov}_{\mathbb{X}}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}])\right)\end{aligned}$$

Decomposing Bias and Variance

- C_{CV} is an estimator of the expected generalization error $E_{\mathbb{X}}(C[f_N])$:

$$\text{MSE}(C_{CV}) = \text{Var}_{\mathbb{X}}(C_{CV}) + \text{bias}(C_{CV})^2, \text{ where}$$

$$\begin{aligned}\text{bias}(C_{CV}) &= E_{\mathbb{X}}(C_{CV}) - E_{\mathbb{X}}(C[f_N]) = E\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) - E(C[f_N]) \\ &= \frac{1}{K} \sum_i E(C[f_{-N^{(i)}}]) - E(C[f_N]) \\ &= E(C[f_{-N^{(s)}}]) - E(C[f_N]), \forall s \\ &= \text{bias}(C[f_{-N^{(s)}}]), \forall s\end{aligned}$$

$$\begin{aligned}\text{Var}_{\mathbb{X}}(C_{CV}) &= \text{Var}\left(\sum_i \frac{1}{K} C[f_{-N^{(i)}}]\right) = \frac{1}{K^2} \text{Var}\left(\sum_i C[f_{-N^{(i)}}]\right) \\ &= \frac{1}{K^2} \left(\sum_i \text{Var}(C[f_{-N^{(i)}}]) + 2 \sum_{i,j,j>i} \text{Cov}_{\mathbb{X}}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}])\right) \\ &= \frac{1}{K} \text{Var}(C[f_{-N^{(s)}}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}]), \forall s\end{aligned}$$

How Many Folds K ? II

$$\text{MSE}(C_{\text{CV}}) = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2, \text{ where}$$

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N^{(s)}}]), \forall s$$

$$\text{Var}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N^{(s)}}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}]), \forall s$$

trade off

- We can reduce $\text{bias}(C_{\text{CV}})$ and $\text{Var}(C_{\text{CV}})$ by *learning theory*
 - Choosing the right model complexity avoiding both underfitting and overfitting
 - Collecting more training examples (N)

How Many Folds K ? II

$$\text{MSE}(C_{\text{CV}}) = \text{Var}_{\mathbb{X}}(C_{\text{CV}}) + \text{bias}(C_{\text{CV}})^2, \text{ where}$$

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N(s)}]), \forall s$$

$$\text{Var}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$

- We can reduce $\text{bias}(C_{\text{CV}})$ and $\text{Var}(C_{\text{CV}})$ by *learning theory*
 - Choosing the right model complexity avoiding both underfitting and overfitting
 - Collecting more training examples (N)
- Furthermore, we can reduce $\text{Var}(C_{\text{CV}})$ by *making $f_{-N(i)}$ and $f_{-N(j)}$ uncorrelated*

How Many Folds K ? III


$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N(s)}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$

- With a large K , the C_{CV} tends to have:

How Many Folds K ? III

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N(s)}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$


- With a large K , the C_{CV} tends to have:
 - Low $\text{bias}(C[f_{-N(s)}])$ and $\text{Var}(C[f_{-N(s)}])$, as $f_{-N(s)}$ is trained on more examples



How Many Folds K ? III

$$\text{bias}(C_{CV}) = \text{bias}(C[f_{-N(s)}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{CV}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$

- With a large K , the C_{CV} tends to have:
 - Low $\text{bias}(C[f_{-N(s)}])$ and $\text{Var}(C[f_{-N(s)}])$, as $f_{-N(s)}$ is trained on more examples
 - High $\text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}])$, as training sets $\mathbb{X} \setminus \mathbb{X}^{(i)}$ and $\mathbb{X} \setminus \mathbb{X}^{(j)}$ are more similar thus $C[f_{-N(i)}]$ and $C[f_{-N(j)}]$ are more positively correlated




K 越大 data set 會被切越細
每次都 99% train 1% validation
那每一次拿來 train 的資料都會很相似
 $\text{Cov}(i,j)$ 會高

How Many Folds K ? IV

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N^{(s)}}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N^{(s)}}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}]), \forall s$$

- Conversely, with a small K , the cross-validation error tends to have a high $\text{bias}(C[f_{-N^{(s)}}])$ and $\text{Var}(C[f_{-N^{(s)}}])$ but low $\text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}])$




How Many Folds K ? IV

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N^{(s)}}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N^{(s)}}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}]), \forall s$$

- Conversely, with a small K , the cross-validation error tends to have a high $\text{bias}(C[f_{-N^{(s)}}])$ and $\text{Var}(C[f_{-N^{(s)}}])$ but low $\text{Cov}(C[f_{-N^{(i)}}], C[f_{-N^{(j)}}])$
- In practice, we usually set $K = 5$ or 10


data set 很大的話選3 or 5 data set 普通的話 選10 可以降低一些var的影響



Leave-One-Out CV

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N(s)}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$


- For very small dataset:
 - $\text{MSE}(C_{\text{CV}})$ is dominated by $\text{bias}(C[f_{-N(s)}])$ and $\text{Var}(C[f_{-N(s)}])$
 - Not $\text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}])$



Leave-One-Out CV

$$\text{bias}(C_{\text{CV}}) = \text{bias}(C[f_{-N(s)}]), \forall s$$
$$\text{Var}_{\mathbb{X}}(C_{\text{CV}}) = \frac{1}{K} \text{Var}(C[f_{-N(s)}]) + \frac{2}{K^2} \sum_{i,j>i} \text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}]), \forall s$$

- For very small dataset:
 - $\text{MSE}(C_{\text{CV}})$ is dominated by $\text{bias}(C[f_{-N(s)}])$ and $\text{Var}(C[f_{-N(s)}])$
 - Not $\text{Cov}(C[f_{-N(i)}], C[f_{-N(j)}])$
- We can choose $K = N$, which we call the *leave-one-out CV*



Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Ensemble Methods

- *No free lunch theorem*: there is no single ML algorithm that always outperforms the others in all domains/tasks

Ensemble Methods

- *No free lunch theorem*: there is no single ML algorithm that always outperforms the others in all domains/tasks
- Can we combine multiple base-learners to improve
 - Applicability across different domains, and/or
 - Generalization performance in a specific task?

Ensemble Methods

- *No free lunch theorem*: there is no single ML algorithm that always outperforms the others in all domains/tasks
- Can we combine multiple base-learners to improve
 - Applicability across different domains, and/or
 - Generalization performance in a specific task?
- These are the goals of *ensemble learning*

Ensemble Methods

- *No free lunch theorem*: there is no single ML algorithm that always outperforms the others in all domains/tasks
- Can we combine multiple base-learners to improve
 - Applicability across different domains, and/or
 - Generalization performance in a specific task?
- These are the goals of *ensemble learning*
- How to “combine” multiple base-learners?

Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Voting

- **Voting**: linear combining the predictions of base-learners for each x :

$$\tilde{y}_k = \sum_j w_j \hat{y}_k^{(j)} \text{ where } w_j \geq 0, \sum_j w_j = 1.$$

- If all learners are given equal weight $w_j = 1/L$, we have the **plurality vote** (multi-class version of majority vote)

Voting Rule	Formular
Sum	$\tilde{y}_k = \frac{1}{L} \sum_{j=1}^L \hat{y}_k^{(j)}$
Weighted sum	$\tilde{y}_k = \sum_j w_j \hat{y}_k^{(j)}, w_j \geq 0, \sum_j w_j = 1$
Median	$\tilde{y}_k = \text{median}_j \hat{y}_k^{(j)}$
Minimum	$\tilde{y}_k = \min_j \hat{y}_k^{(j)}$
Maximum	$\tilde{y}_k = \max_j \hat{y}_k^{(j)}$
Product	$\tilde{y}_k = \prod_j \hat{y}_k^{(j)}$

Why Voting Works? I

Why Voting Works? I

- Assume that each $\hat{y}^{(j)}$ has the expected value $E_{\mathbb{X}}(\hat{y}^{(j)} | \mathbf{x})$ and variance $\text{Var}_{\mathbb{X}}(\hat{y}^{(j)} | \mathbf{x})$
- When $w_j = 1/L$, we have:

$$E_{\mathbb{X}}(\tilde{y} | \mathbf{x}) = E\left(\sum_j \frac{1}{L} \hat{y}^{(j)} | \mathbf{x}\right) = \frac{1}{L} \sum_j E(\hat{y}^{(j)} | \mathbf{x}) = E(\hat{y}^{(j)} | \mathbf{x})$$

$$\begin{aligned}\text{Var}_{\mathbb{X}}(\tilde{y} | \mathbf{x}) &= \text{Var}\left(\sum_j \frac{1}{L} \hat{y}^{(j)} | \mathbf{x}\right) = \frac{1}{L^2} \text{Var}\left(\sum_j \hat{y}^{(j)} | \mathbf{x}\right) \\ &= \frac{1}{L} \text{Var}(\hat{y}^{(j)} | \mathbf{x}) + \frac{2}{L^2} \sum_{i,j, i < j} \text{Cov}(\hat{y}^{(i)}, \hat{y}^{(j)} | \mathbf{x})\end{aligned}$$

Why Voting Works? I

- Assume that each $\hat{y}^{(j)}$ has the expected value $E_{\mathbb{X}}(\hat{y}^{(j)} | \mathbf{x})$ and variance $\text{Var}_{\mathbb{X}}(\hat{y}^{(j)} | \mathbf{x})$
- When $w_j = 1/L$, we have:

$$E_{\mathbb{X}}(\tilde{y} | \mathbf{x}) = E\left(\sum_j \frac{1}{L} \hat{y}^{(j)} | \mathbf{x}\right) = \frac{1}{L} \sum_j E(\hat{y}^{(j)} | \mathbf{x}) = E(\hat{y}^{(j)} | \mathbf{x})$$

$$\begin{aligned}\text{Var}_{\mathbb{X}}(\tilde{y} | \mathbf{x}) &= \text{Var}\left(\sum_j \frac{1}{L} \hat{y}^{(j)} | \mathbf{x}\right) = \frac{1}{L^2} \text{Var}\left(\sum_j \hat{y}^{(j)} | \mathbf{x}\right) \\ &= \frac{1}{L} \text{Var}(\hat{y}^{(j)} | \mathbf{x}) + \frac{2}{L^2} \sum_{i,j, i < j} \text{Cov}(\hat{y}^{(i)}, \hat{y}^{(j)} | \mathbf{x})\end{aligned}$$

- The expected value doesn't change, so the bias doesn't change

Why Voting Works? II

$$\text{Var}_{\mathbb{X}}(\tilde{y}|\mathbf{x}) = \frac{1}{L} \text{Var}\left(\hat{y}^{(j)}|\mathbf{x}\right) + \frac{2}{L^2} \sum_{i,j,i < j} \text{Cov}\left(\hat{y}^{(i)}, \hat{y}^{(j)}|\mathbf{x}\right)$$

Why Voting Works? II

$$\text{Var}_{\mathbb{X}}(\tilde{y}|\mathbf{x}) = \frac{1}{L} \text{Var}(\hat{y}^{(j)}|\mathbf{x}) + \frac{2}{L^2} \sum_{i,j,i < j} \text{Cov}(\hat{y}^{(i)}, \hat{y}^{(j)}|\mathbf{x})$$

- If $\hat{y}^{(i)}$ and $\hat{y}^{(j)}$ are uncorrelated, the variance can be reduced

Why Voting Works? II

$$\text{Var}_{\mathbb{X}}(\tilde{y}|\mathbf{x}) = \frac{1}{L} \text{Var}(\hat{y}^{(j)}|\mathbf{x}) + \frac{2}{L^2} \sum_{i,j,i < j} \text{Cov}(\hat{y}^{(i)}, \hat{y}^{(j)}|\mathbf{x})$$

- If $\hat{y}^{(i)}$ and $\hat{y}^{(j)}$ are uncorrelated, the variance can be reduced
- Unfortunately, $\hat{y}^{(j)}$'s may **not** be i.i.d. in practice
- If voters are positively correlated, variance increases

Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Bagging

- *Bagging* (short for *bootstrap aggregating*) is a voting method, but base-learners are made different deliberately
- How?

Bagging

- *Bagging* (short for *bootstrap aggregating*) is a voting method, but base-learners are made different deliberately
- How? Why not train them using slightly different training sets?

Bagging

- **Bagging** (short for **bootstrap aggregating**) is a voting method, but base-learners are made different deliberately
- How? Why not train them using slightly different training sets?
- ① Generate L slightly different samples from a given sample is done by **bootstrap**: given \mathbb{X} of size N , we draw N points randomly from \mathbb{X} **with replacement** to get $\mathbb{X}^{(j)}$
 - It is possible that some instances are drawn more than once and some are not at all

Bagging

- **Bagging** (short for **bootstrap aggregating**) is a voting method, but base-learners are made different deliberately
- How? Why not train them using slightly different training sets?
- ① Generate L slightly different samples from a given sample is done by **bootstrap**: given \mathbb{X} of size N , we draw N points randomly from \mathbb{X} **with replacement** to get $\mathbb{X}^{(j)}$
 - It is possible that some instances are drawn more than once and some are not at all
- ② Train a base-learner for each $\mathbb{X}^{(j)}$

Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Boosting

- In bagging, generating “uncorrelated” base-learners is left to chance and instability of the learning method

Boosting

- In bagging, generating “uncorrelated” base-learners is left to chance and instability of the learning method
- In *boosting*, we generate *complementary* base-learners
- How?

Boosting

- In bagging, generating “uncorrelated” base-learners is left to chance and instability of the learning method
- In *boosting*, we generate *complementary* base-learners
- How? Why not train the next learner on the mistakes of the previous learners

Boosting

- In bagging, generating “uncorrelated” base-learners is left to chance and instability of the learning method
- In *boosting*, we generate *complementary* base-learners
- How? Why not train the next learner on the mistakes of the previous learners
- For simplicity, let's consider the binary classification here:
 $d^{(j)}(\mathbf{x}) \in \{1, -1\}$
- The original boosting algorithm combines three *weak learners* to generate a *strong learner*
 - A weak learner has error probability less than $1/2$ (better than random guessing)
 - A strong learner has arbitrarily small error probability

Original Boosting Algorithm

- Training
- ① Given a large training set, randomly divide it into three

Original Boosting Algorithm

- Training
- ① Given a large training set, randomly divide it into three
 - ② Use $\mathbb{X}^{(1)}$ to train the first learner $d^{(1)}$ and feed $\mathbb{X}^{(2)}$ to $d^{(1)}$

Original Boosting Algorithm

- Training
 - ① Given a large training set, randomly divide it into three
 - ② Use $\mathbb{X}^{(1)}$ to train the first learner $d^{(1)}$ and feed $\mathbb{X}^{(2)}$ to $d^{(1)}$
 - ③ Use all points misclassified by $d^{(1)}$ and $\mathbb{X}^{(2)}$ to train $d^{(2)}$. Then feed $\mathbb{X}^{(3)}$ to $d^{(1)}$ and $d^{(2)}$

Original Boosting Algorithm

- Training
 - ① Given a large training set, randomly divide it into three
 - ② Use $\mathbb{X}^{(1)}$ to train the first learner $d^{(1)}$ and feed $\mathbb{X}^{(2)}$ to $d^{(1)}$
 - ③ Use all points misclassified by $d^{(1)}$ and $\mathbb{X}^{(2)}$ to train $d^{(2)}$. Then feed $\mathbb{X}^{(3)}$ to $d^{(1)}$ and $d^{(2)}$
 - ④ Use the points on which $d^{(1)}$ and $d^{(2)}$ disagree to train $d^{(3)}$

Original Boosting Algorithm

- Training

- ① Given a large training set, randomly divide it into three
- ② Use $\mathbb{X}^{(1)}$ to train the first learner $d^{(1)}$ and feed $\mathbb{X}^{(2)}$ to $d^{(1)}$
- ③ Use all points misclassified by $d^{(1)}$ and $\mathbb{X}^{(2)}$ to train $d^{(2)}$. Then feed $\mathbb{X}^{(3)}$ to $d^{(1)}$ and $d^{(2)}$
- ④ Use the points on which $d^{(1)}$ and $d^{(2)}$ disagree to train $d^{(3)}$

- Testing

- ① Feed a point it to $d^{(1)}$ and $d^{(2)}$ first. If their outputs agree, use them as the final prediction

Original Boosting Algorithm

- Training


- ① Given a large training set, randomly divide it into three
- ② Use $\mathbb{X}^{(1)}$ to train the first learner $d^{(1)}$ and feed $\mathbb{X}^{(2)}$ to $d^{(1)}$
- ③ Use all points misclassified by $d^{(1)}$ and $\mathbb{X}^{(2)}$ to train $d^{(2)}$. Then feed $\mathbb{X}^{(3)}$ to $d^{(1)}$ and $d^{(2)}$
- ④ Use the points on which $d^{(1)}$ and $d^{(2)}$ disagree to train $d^{(3)}$
d1和d2預測的不同的那些. 丟去train d3

- Testing

- ① Feed a point it to $d^{(1)}$ and $d^{(2)}$ first. If their outputs agree, use them as the final prediction
- ② Otherwise the output of $d^{(3)}$ is taken

Example

- Assuming $\mathbb{X}^{(1)}$, $\mathbb{X}^{(2)}$, and $\mathbb{X}^{(3)}$ are the same:



- Disadvantage: requires a large training set to afford the three-way split

AdaBoost

- *AdaBoost*: uses the same training set over and over again
- How to make some points “larger?”

AdaBoost

- *AdaBoost*: uses the same training set over and over again
- How to make some points “larger?”
- Modify the probabilities of drawing the instances as a function of error

AdaBoost

- **AdaBoost**: uses the same training set over and over again
- How to make some points “larger?”
- Modify the probabilities of drawing the instances as a function of error
- Notation:
- $\Pr^{(i,j)}$: probability that an example $(\mathbf{x}^{(i)}, y^{(i)})$ is drawn to train the j th base-learner $d^{(j)}$
- $\epsilon^{(j)} = \sum_i \Pr^{(i,j)} 1(y^{(i)} \neq d^{(j)}(\mathbf{x}^{(i)}))$: error rate of $d^{(j)}$ on its training set

Algorithm

- Training
 - ① Initialize $\Pr^{(i,1)} = \frac{1}{N}$ for all i
 - ② Start from $j = 1$:
 - ① Randomly draw N examples from \mathbb{X} with probabilities $\Pr^{(i,j)}$ and use them to train $d^{(j)}$

Algorithm

- Training
 - ① Initialize $\Pr^{(i,1)} = \frac{1}{N}$ for all i
 - ② Start from $j = 1$:
 - ① Randomly draw N examples from \mathbb{X} with probabilities $\Pr^{(i,j)}$ and use them to train $d^{(j)}$
 - ② Stop adding new base-learners if $\epsilon^{(j)} \geq \frac{1}{2}$

Algorithm

- Training

① Initialize $\Pr^{(i,1)} = \frac{1}{N}$ for all i

② Start from $j = 1$:

① Randomly draw N examples from \mathbb{X} with probabilities $\Pr^{(i,j)}$ and use them to train $d^{(j)}$

② Stop adding new base-learners if $\epsilon^{(j)} \geq \frac{1}{2}$

③ Define $\alpha_j = \frac{1}{2} \log \left(\frac{1 - \epsilon^{(j)}}{\epsilon^{(j)}} \right) > 0$ and set

$$\Pr^{(i,j+1)} = \Pr^{(i,j)} \cdot \exp(-\alpha_j y^{(i)} d^{(j)}(\mathbf{x}^{(i)})) \text{ for all } i$$

預測對的話 $\exp(-a) = 1 / \exp(a)$ 機率變小
預測錯的話 $\exp(a)$ 抽到的機率變大

Algorithm

- Training

- ① Initialize $\Pr^{(i,1)} = \frac{1}{N}$ for all i
- ② Start from $j = 1$:
 - ① Randomly draw N examples from \mathbb{X} with probabilities $\Pr^{(i,j)}$ and use them to train $d^{(j)}$
 - ② Stop adding new base-learners if $\epsilon^{(j)} \geq \frac{1}{2}$
 - ③ Define $\alpha_j = \frac{1}{2} \log \left(\frac{1 - \epsilon^{(j)}}{\epsilon^{(j)}} \right) > 0$ and set $\Pr^{(i,j+1)} = \Pr^{(i,j)} \cdot \exp(-\alpha_j y^{(i)} d^{(j)}(\mathbf{x}^{(i)}))$ for all i
 - ④ Normalize $\Pr^{(i,j+1)}, \forall i$, by multiplying $\left(\sum_i \Pr^{(i,j+1)} \right)^{-1}$

再normalize讓所有機率加起來還是1

Algorithm


- Training

- ① Initialize $\Pr^{(i,1)} = \frac{1}{N}$ for all i
- ② Start from $j = 1$:
 - ① Randomly draw N examples from \mathbb{X} with probabilities $\Pr^{(i,j)}$ and use them to train $d^{(j)}$
 - ② Stop adding new base-learners if $\epsilon^{(j)} \geq \frac{1}{2}$
 - ③ Define $\alpha_j = \frac{1}{2} \log \left(\frac{1 - \epsilon^{(j)}}{\epsilon^{(j)}} \right) > 0$ and set $\Pr^{(i,j+1)} = \Pr^{(i,j)} \cdot \exp(-\alpha_j y^{(i)} d^{(j)}(\mathbf{x}^{(i)}))$ for all i
 - ④ Normalize $\Pr^{(i,j+1)}$, $\forall i$, by multiplying $\left(\sum_i \Pr^{(i,j+1)} \right)^{-1}$

- Testing

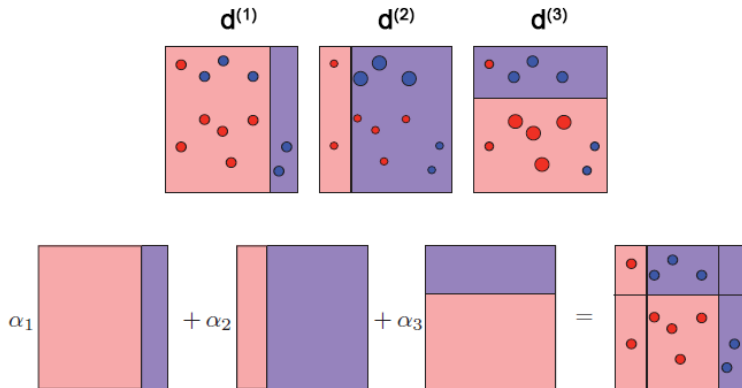
- ① Given \mathbf{x} , calculate $\hat{y}^{(j)}$ for all j
- ② Make final prediction \tilde{y} by voting: $\tilde{y} = \sum_j \alpha_j d^{(j)}(\mathbf{x})$

Example



- $d^{(j+1)}$ complements $d^{(j)}$ and $d^{(j-1)}$ by focusing on predictions they disagree

Example



- $d^{(j+1)}$ complements $d^{(j)}$ and $d^{(j-1)}$ by focusing on predictions they disagree
- Voting weights ($\alpha_j = \frac{1}{2} \log \left(\frac{1 - \epsilon^{(j)}}{\epsilon^{(j)}} \right)$) in predictions are proportional to the base-learner's accuracy

Outline

1 Cross Validation

- How Many Folds?

2 Ensemble Methods

- Voting
- Bagging
- Boosting
- Why AdaBoost Works?

Why AdaBoost Works


- Why AdaBoost improves performance?

Why AdaBoost Works

- Why AdaBoost improves performance?
- By increasing model complexity?

Why AdaBoost Works


- Why AdaBoost improves performance?
- By increasing model complexity? Not exactly
 - Empirical study: AdaBoost **reduces overfitting** as L grows, even when there is no training error



C4.5 decision trees (Schapire et al., 1998).

Why AdaBoost Works


- Why AdaBoost improves performance?
- By increasing model complexity? Not exactly
 - Empirical study: AdaBoost *reduces overfitting* as L grows, even when there is no training error
- AdaBoost *increases margin* [1, 2]



C4.5 decision trees (Schapire et al., 1998).


Margin as Confidence of Predictions

- Recall in SVC, a larger margin improves generalizability




Margin as Confidence of Predictions

- Recall in SVC, a larger margin improves generalizability
- Due to *higher confidence predictions* over training examples



Margin as Confidence of Predictions

- Recall in SVC, a larger margin improves generalizability
- Due to *higher confidence predictions* over training examples




- We can define the margin for AdaBoost similarly
- In binary classification, define *margin* of a prediction of an example $(\mathbf{x}^{(i)}, y^{(i)}) \in \mathbb{X}$ as:

$$\text{margin}(\mathbf{x}^{(i)}, y^{(i)}) = y^{(i)} f(\mathbf{x}^{(i)}) = \sum_{j: y^{(i)} = d^{(j)}(\mathbf{x}^{(i)})} \alpha_j - \sum_{j: y^{(i)} \neq d^{(j)}(\mathbf{x}^{(i)})} \alpha_j$$

Margin Distribution

- Margin distribution over θ :

$$\Pr_{\mathbb{X}}(y^{(i)}f(\mathbf{x}^{(i)}) \leq \theta) \approx \frac{|(\mathbf{x}^{(i)}, y^{(i)}) : y^{(i)}f(\mathbf{x}^{(i)}) \leq \theta|}{|\mathbb{X}|}$$




increase rounds 大家的margin都提高了

Margin Distribution


- Margin distribution over θ :

$$\Pr_{\mathbb{X}}(y^{(i)}f(\mathbf{x}^{(i)}) \leq \theta) \approx \frac{|\{(\mathbf{x}^{(i)}, y^{(i)}) : y^{(i)}f(\mathbf{x}^{(i)}) \leq \theta\}|}{|\mathbb{X}|}$$



LEGEND: (small dash, large dash, solid) lines equal (5, 100, 1000) rounds of boosting

- A complementary learner:
- Clarifies low confidence areas
- Increases margin of points in these areas



Reference I

- [1] Yoav Freund, Robert Schapire, and N Abe.
A short introduction to boosting.
Journal-Japanese Society For Artificial Intelligence, 14(771-780):1612,
1999.
- [2] Liwei Wang, Masashi Sugiyama, Cheng Yang, Zhi-Hua Zhou, and Jufu Feng.
On the margin explanation of boosting algorithms.
In *COLT*, pages 479–490. Citeseer, 2008.